

Hidden Instruction Authority Cross

Visible Interaction, Invisible Governance, and Authority Asymmetry in AI Systems

Version: 0.1

Date: July 2026

Status: Preprint / OSF external-facing draft

Context: Meta-Writing Ecology

Document Type: Cross

Huang Tzuyuan

Abstract

This document formalizes the **Hidden Instruction Authority Cross** as a structural intersection between visible interaction and invisible high-authority instruction. The Cross appears when a user can operate an AI interface, provide prompts, select options, and receive outputs while remaining unable to inspect the instruction hierarchy that materially shapes the system's behavior.

The condition is not reducible to secrecy or poor documentation. It arises from the combination of an accessible interaction surface, a higher-authority instruction layer, limited user control, and responsibility for interpreting the resulting behavior. System-level instructions, developer directives, policy constraints, safety rules, persona settings, tool boundaries, memory policies, and deployment configurations may have greater authority than user input while remaining partially or entirely unavailable to the user.

This produces a control–visibility gap. Users may misattribute refusals, styles, biases, omissions, or behavioral changes to model capability, model personality, or their own prompting when the immediate cause lies in a hidden governance layer. Surface customization can intensify the confusion by creating an appearance of control without access to the underlying authority structure.

The Cross does not demand disclosure of every system instruction or the removal of safety, security, or proprietary boundaries. It provides a framework for distinguishing operational access from governance visibility and for examining how explanation, contestability, versioning, and responsibility should be designed when the governing instruction stack cannot be fully exposed.

Keywords

Hidden Instructions; Instruction Authority; System Prompt; Developer Instructions; AI Governance; Visible Interaction; Authority Asymmetry; Control–Visibility Gap; Misattribution; User Agency; Transparency; Customization; Boundary Integration; Meta-Writing Ecology.

1. Introduction

Conversational interfaces present a direct relation: the user writes, the system responds. This surface can make system behavior appear to arise from the model and the user alone. In deployed systems, the interaction is usually conditioned by additional layers that the user did not author and may not be able to inspect.

These layers are not incidental. They can define the task, prioritize instructions, prohibit actions, shape tone, control tools, insert context, and determine when refusal occurs. The user remains operationally active but constitutionally subordinate to an unseen instruction hierarchy.

The resulting problem is not solved by giving the user more prompting tips. It requires a distinction between the capacity to operate an interface and visibility into the authority that governs it. This Cross examines that intersection and the misattribution, burden, and accountability problems it can produce.

2. Core Definition

Hidden Instruction Authority Cross names a cross-condition in which a user enters a visible interaction surface while the actual conditions of interaction have

already been configured by a higher-authority instruction layer that the user cannot fully see, inspect, contest, modify, or audit.

In LLM systems, this layer appears as the system prompt. The system prompt guides behaviour, defines interaction boundaries, and may override user input. However, the user often encounters only the visible conversational interface, not the higher-order instruction structure that governs what the system can say, refuse, prioritize, obscure, or reframe.

The core issue is not whether users can enter prompts.

The core issue is:

Can users see, understand, trace, question, or modify the higher-authority layer that preconditions the interaction?

When they cannot, users receive operational access without structural visibility.

3. Classification and Scope

This document presents **Hidden Instruction Authority Cross** as a Cross. A Cross identifies a stable intersection among existing mechanisms, roles, or system conditions. Its function is to preserve the distinct source structures while making their interaction analyzable. It should not be converted into a standalone model merely because the intersection is recurrent.

The Cross is external-facing and conceptual. It does not establish empirical prevalence, legal duty, technical compliance, or a universal evaluation rule. The public version omits non-public editorial records, private source materials, non-public working materials, and implementation details.

4. Central Distinctions

Visible interaction $\neq$ visible governance Prompting access $\neq$ instruction authority
Customization $\neq$ control Disclosure $\neq$ accountability Control option $\neq$ responsibility
alignment System behaviour $\neq$ model capability alone

These distinctions are necessary because a user can appear to control an interaction while the interaction remains governed by an unseen authority layer.

5. Core Mechanism and Structural Sequence

5.1 Core Mechanism

The Cross is produced by the separation of three layers:

1. Visible interaction layer
2. Hidden instruction-authority layer
3. Responsibility and interpretation layer

The visible interaction layer gives the user a sense of direct access.

The hidden instruction-authority layer defines what counts as permissible, preferred, safe, stylistically appropriate, or structurally available.

The responsibility and interpretation layer determines who must explain, absorb, correct, or trust the resulting behaviour.

The mechanism activates when these three layers do not align.

5.2 Structural Sequence

User enters visible interface ↓ User input appears to initiate the interaction ↓
Hidden high-authority instruction layer preconditions the response field ↓ Model

output appears as a direct response to the user ↓ User interprets the output as model capability, personality, bias, refusal, limitation, or judgment ↓ Actual cause may lie in hidden governance instructions ↓ Interpretive burden shifts toward the user-facing surface ↓ Responsibility, trust, control, and accountability become misaligned

6. Structural Components

1. Visible Interaction Surface

The user sees an interface that appears open, responsive, and conversational.

This surface may include:

chat input model output settings persona options style controls custom
instructions tool access model choice

The surface creates an impression of user-driven interaction.

However, interaction visibility does not guarantee structural legibility.

2. Hidden High-Authority Instruction Layer

The system includes a higher-authority instruction layer that preconditions the interaction.

In LLM systems, this may include:

system prompts developer messages platform policies safety instructions style
constraints refusal rules persona definitions memory policies tool-use boundaries
legal restrictions deployment-specific directives

The user may not know which instruction is active, which layer has priority, or which actor authored the constraint.

3. Authority Asymmetry

The instruction layer has more authority than the user's input.

The user can request, but the higher-authority layer can override.

This creates an asymmetry:

The user participates in the interaction. The system authority configures the interaction.

The user's role is operational.

The hidden instruction layer's role is constitutional.

4. Misattribution Field

When the user cannot see the governing instruction layer, system behaviour may be misattributed.

A user may interpret output as:

model intelligence model limitation model bias model personality technical failure ethical judgment platform position user error

But the actual cause may be:

system prompt constraint developer instruction policy layer safety rule style rule hidden refusal condition retrieval or tool restriction deployment configuration

This creates a misattribution field.

5. Control–Visibility Gap

The user may be given limited control mechanisms, but not visibility into the deeper authority structure.

Examples:

style settings without system prompt access custom instructions without platform-level visibility persona controls without safety-layer transparency model

selection without instruction-stack disclosure settings access without versioned prompt history

This produces a gap between what the user can change and what actually governs the interaction.

6. Responsibility Displacement

If users are given surface-level control without clear responsibility boundaries, interpretive and corrective labor may shift toward them.

The user may be expected to:

guess why the system refused test alternative prompts infer hidden constraints work around unexplained behaviour judge whether output reflects policy or capability detect bias without seeing its source maintain trust without sufficient visibility

This is not full agency. It is burden relocation.

7. Primary Failure Patterns

1. Interaction-Surface Overreach

The openness of the user interface is interpreted as visibility into the system's governing conditions.

2. Model-Policy Misattribution

Behavior produced by deployment rules or higher-authority instructions is attributed to model capability, personality, or knowledge.

3. Customization Illusion

User-facing settings create an appearance of control while the decisive instruction layers remain unavailable.

4. Hidden Version Drift

Instruction or policy layers change without a visible record, causing unexplained behavioral shifts.

5. User-Side Debugging Burden

Users are required to infer hidden constraints through repeated prompting, trial, and error.

6. Responsibility Without Authority

Users are treated as responsible operators even though the governing conditions are authored and controlled elsewhere.

8. Operational Evaluation

Use the following questions to identify the Cross:

1. Is the user interacting through a visible interface?
2. Is there a higher-authority instruction layer above user input?
3. Can the user see that layer?
4. Can the user identify which instruction caused a behaviour?
5. Can the user distinguish policy from capability?
6. Can the user distinguish safety from style?
7. Can the user distinguish refusal from technical limitation?
8. Can the user inspect version changes?
9. Can the user contest or configure the relevant boundary?

10. Does the system shift interpretive or corrective labor to the user?

If most answers indicate hidden authority plus user-side interpretation burden, the Cross is active.

Failure Pattern

The user receives output. The output appears to express the model. The model appears to be the source. The hidden instruction layer remains absent. The user interprets the behaviour without the governing condition. Misattribution increases. Trust, agency, and accountability destabilize.

9. Application Domains

1. Conversational AI

Users may experience unexplained refusals or behavioral shifts without knowing which authority layer caused them.

2. Enterprise and institutional deployments

Organization-specific instructions can shape outputs while employees or clients see only the general interface.

3. Tool-using assistants

Hidden rules can govern when tools are available, what actions are permitted, and which information is withheld.

4. Personalization and memory

User-facing controls may coexist with provider-controlled memory, policy, or profile layers that have greater authority.

5. AI-mediated public information

Outputs shaped by hidden instructions may circulate as if they were neutral statements of model capability or knowledge.

10. Relations to Adjacent Public Frameworks

1. False Legibility Model

False Legibility Model describes how readable surfaces can be accepted before their underlying structure is understood. The present entry identifies a concrete way in which fluency, familiarity, or interface clarity can conceal an unresolved boundary.

2. Responsibility Alignment Model

Responsibility Alignment Model clarifies whether ownership, authority, capacity, visibility, and accountability remain attached to compatible nodes. The present entry specifies one condition under which that alignment can fail or become nominal.

3. Boundary Integration Failure

Boundary Integration Failure describes collapse produced by merging roles, instructions, or authority layers that require differentiation. The present entry identifies one layer pair that must remain distinguishable.

4. Cost Visibility and Redistribution Model

Cost Visibility and Redistribution Model tracks how costs become visible at one node and hidden or displaced at another. The present entry shows how verification, correction, access, or responsibility costs can be redistributed by the relevant interface.

5. AI-Readable Knowledge Architecture

AI-Readable Knowledge Architecture concerns the structuring of public knowledge so that machine-mediated reading does not erase provenance, limits, or relation status. The present entry supplies a boundary that such architecture should keep visible.

11. Boundary Conditions and Non-Applicability

11.1 Boundary Conditions

This Cross does not require full public disclosure of all system instructions.

It requires structural recognition that:

hidden authority exists hidden authority shapes output hidden authority affects user interpretation hidden authority may relocate burden hidden authority requires appropriate disclosure, auditing, or accountability

A system can preserve necessary safety or security boundaries while still documenting the existence, role, scope, versioning, and accountability path of hidden instruction layers.

Relation to Transparency

Transparency in this Cross does not mean total exposure.

It means sufficient visibility into the conditions that govern interaction.

Possible transparency forms include:

summary-level prompt disclosure topic-level instruction disclosure behaviour-triggered explanation version-change notices prompt cards settings-page

documentation authority-layer maps refusal-source indicators audit logs stakeholder-specific disclosure

The relevant question is not:

Should everything be visible?

The relevant question is:

What must be visible for users to understand the governing conditions of the interaction?

Relation to Control

Control in this Cross does not mean unrestricted editing.

Unrestricted control may create safety, liability, or comprehension problems.

Structurally safer control forms include:

pre-set modes preference configuration task-specific settings session-level configuration bounded user instructions disclosure before system updates settings-based review controlled override pathways human-readable explanation of limits

The relevant distinction is:

Agency through structured control $\neq$ liability transfer through unrestricted editing

Relation to Accountability

Accountability requires that system instructions be treated as governance artifacts rather than merely technical configuration.

System prompts and related instruction layers should therefore be understood as:

versioned auditable authored justified bounded change-tracked stakeholder-relevant impact-bearing

For governance purposes, material instruction layers can be treated as versioned, auditable artifacts that document what changed, when, why, and under whose authority.

Minimal Example

A user asks a model a question. The model refuses. The user interprets the refusal as model limitation or moral judgment. The actual cause is a hidden system instruction. The user cannot see the instruction, identify its source, inspect its version, or know whether it comes from safety policy, legal risk, platform preference, or deployment context. The user has access to the interface but not to the governing condition.

This is the Hidden Instruction Authority Cross.

11.2 Non-Applicability

The framework should not be used as a universal criticism of the relevant technology, workflow, role, or interface. It applies to a specific structural mismatch, not to every instance of the general category.

12. Public Orientation and Method Boundary

This document is designed as a public conceptual entry. It provides a definition, mechanism, diagnostic orientation, and limit statement. It omits non-public editorial records, private source materials, non-public working materials, and implementation details.

The entry is not a substitute for empirical research, legal analysis, professional judgment, technical testing, or domain-specific governance. It should not be used to score people or institutions, automate high-stakes decisions, or infer misconduct from conceptual similarity.

Examples and application domains are illustrative. A case should be classified under **Hidden Instruction Authority Cross** only when the relevant structure can be demonstrated from available evidence.

13. Limitations

1. Legitimate non-disclosure

Security, abuse prevention, privacy, and proprietary constraints may justify withholding some instructions.

2. No demand for raw prompt publication

Meaningful transparency can use categories, change logs, explanation surfaces, and appeal mechanisms rather than full disclosure.

3. Model behavior has multiple causes

Not every output can be attributed to a specific instruction layer.

4. User control is necessarily bounded

The Cross does not imply that users should have constitutional authority over deployed systems.

5. No legal transparency standard

The framework does not determine regulatory duties.

14. Minimal Formulation

When a visible interaction surface is governed by an invisible high-authority instruction layer, the user receives operational access without structural visibility into the conditions that shape the interaction.

Alternative compressed form:

Prompting access is not governance visibility.

Expanded form:

A user may be able to operate an interface without being able to inspect the authority structure that determines the interface's behavioural limits, priorities, refusals, and interpretive outputs.

15. Conclusion

Prompting access is not governance visibility. A user may participate extensively in an interaction while remaining unable to see the higher-authority conditions that shape what the system can say or do.

The Hidden Instruction Authority Cross makes this asymmetry analyzable without requiring total exposure. It supports designs that distinguish model capability from deployment policy, explain material boundaries, preserve versioned governance information, and avoid assigning users responsibility for behavior governed elsewhere.

Naming Declaration

Hidden Instruction Authority Cross names a structural intersection within Meta-writing Ecology.

It refers to the condition in which visible interaction is governed by a hidden or low-visibility high-authority instruction layer, producing a gap between user operation, system governance, interpretive visibility, and responsibility allocation.

This term does not refer only to system prompts.

System prompts are the current technical case through which the structure becomes visible.

The Cross may apply to any system where:

1. A user interacts through a visible surface.
2. A higher-authority layer preconditions what can occur.
3. The higher-authority layer is hidden, inaccessible, or only partially disclosed.
4. User input is structurally subordinate to that layer.
5. Outputs are interpreted as direct system behaviour.
6. The governing instruction source cannot be fully traced by the user.
7. Control, trust, and responsibility become misaligned.

The term does not imply that all hidden instructions are illegitimate.

It identifies the structural risk produced when authority, visibility, control, and accountability are not aligned.

Operational Definition

A Hidden Instruction Authority Cross is present when:

A visible interaction surface gives users operational access, but a higher-authority instruction layer determines the conditions of interaction without equivalent user visibility, traceability, or accountability.